

CIT Complex IT Systems

Section 1

Database Programming

Troels Andreassen

Outline

- ❑ Accessing SQL From a Programming Language, using API's, Exemplified
 - JDBC, ODBC and ADO.NET
 - ADO.NET with C# will be covered in more detail in CIT Section 2
- ❑ Programming the database
 - Functions and Procedural Constructs in SQL
 - Triggers in SQL
- ❑ Recursion in SQL

Accessing SQL from a Programming Language

A database programmer must have access to a general-purpose programming language for at least two reasons

- ❑ Not all queries can be expressed in SQL, since SQL does not provide the full expressive power of a general-purpose language.
- ❑ Non-declarative actions -- such as printing a report, interacting with a user, or sending the results of a query to a graphical user interface -- cannot be done from within SQL.

JDBC and ODBC and ADO.NET

- ❑ API's (application-program interfaces) for a program to interact with a database server
- ❑ Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- ❑ JDBC (Java Database Connectivity)
 - works with Java
- ❑ ODBC (Open Database Connectivity)
 - works with C, C++, C#, and Visual Basic
- ❑ ADO.NET
 - works with the .NET framework
 - will be used with C# on the CIT course
 - In particular we will consider what's called Object Relational Mapping (ORM) with Entity-Framework and using LINQ in C#
 - More about this in section 2

JDBC

- ❑ **JDBC** is a Java API for communicating with database systems supporting SQL.
- ❑ JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- ❑ Approach for communicating with the database:
 - 1) Open a connection
 - 2) Create a “statement” object
 - 3) **Execute query**
 - 4) Extract data from result set
 - 5) Close connection
 - (Use exception mechanism to handle errors)

```
//STEP 1. Import required packages
```

```
import java.sql.*;
```

Java & JDBC Code

```
public class FirstExample {
```

```
    static final String DB_URL = "jdbc:postgresql://localhost:5432/university"; // a JDBC url
```

```
    //static final String DB_URL
```

```
    static final String USER = "postgres";
```

```
    static final String PASS = "toor";
```

```
    public static void main(String[] args) {
```

```
        Connection conn = null;
```

```
        Statement stmt = null;
```

```
        try{
```

```
            Class.forName("org.postgresql.Driver");
```

// 1) Open a connection

```
            System.out.println("Connecting to database...");
```

```
            conn = DriverManager.getConnection(DB_URL,USER,PASS);
```

// 2) Create a "statement" object

```
            System.out.println("Creating statement...");
```

```
            stmt = conn.createStatement();
```

```
            String sql;
```

```
            sql = "SELECT id, name, salary FROM instructor";
```

// 3) Execute query

```
            ResultSet rs = stmt.executeQuery(sql);
```

// 4) Extract data from result set

```
            while(rs.next()){
```

```
                //Retrieve by column name and display values
```

```
                System.out.println("ID: " + rs.getString("id") + " " + rs.getString("name") + " " + rs.getInt("salary"));
```

```
            }
```

Connecting to database...

Creating statement...

ID: 10101 Srinivasan 65000

ID: 12121 Wu 90000

ID: 15151 Mozart 40000

ID: 22222 Einstein 95000

ID: 32343 El Said 60000

ID: 33456 Gold 87000

ID: 45565 Katz 75000

ID: 58583 Califiori 62000

ID: 76543 Singh 80000

ID: 76766 Crick 72000

ID: 83821 Brandt 92000

ID: 98345 Kim 80000

Goodbye!

// 5) Close connection

```
            rs.close();
```

```
            stmt.close();
```

```
            conn.close();
```

```
        }catch(Exception e){
```

```
            //Handle errors
```

```
            e.printStackTrace();
```

```
        }
```

```
        System.out.println("Goodbye!");
```

```
    }//end main
```

```
}//end FirstExample
```

ADO.NET

- ❑ The ADO.NET API provides functions to access data similar to the JDBC functions.
- ❑ Thus ADO.NET allows access to results of SQL queries
- ❑ Overall the approach for communicating with the database is the same:
 - 1) Open a connection
 - 2) Create a “statement” object
 - 3) Execute query
 - 4) Extract data from result set
 - 5) Close connection

```
using System;
using Npgsql;
```

C# & ADO.NET

```
namespace AdoExample
```

```
{
```

```
    class Program
```

```
    {
```

```
        static void Main(string[] args)
```

```
        {
```

```
            var connString = "Host=localhost;Username=postgres;Password=xxxx;Database=university";
```

```
            using (var conn = new NpgsqlConnection(connString))
```

```
            {
```

```
                // 1) Open a connection
```

```
                conn.Open();
```

```
                // 2) Create a "statement" object
```

```
                using (var cmd = new NpgsqlCommand("SELECT id, name, salary FROM instructor", conn))
```

```
                // 3) Execute query
```

```
                using (var rdr = cmd.ExecuteReader())
```

```
                // 4) Extract data from result set
```

```
                while (rdr.Read()){
```

```
                    Console.WriteLine("ID: {0} {1} {2}", rdr[0], rdr[1], rdr[2]);
```

```
                }
```

```
                // 5) Close the connection
```

```
                conn.Close();
```

```
            }
```

```
        }
```

```
    }
```

```
}
```

ID: 10101 Srinivasan 65000.00

ID: 12121 Wu 90000.00

ID: 15151 Mozart 40000.00

ID: 22222 Einstein 95000.00

ID: 32343 El Said 60000.00

ID: 33456 Gold 87000.00

ID: 45565 Katz 75000.00

ID: 58583 Califieri 62000.00

ID: 76543 Singh 80000.00

ID: 76766 Crick 72000.00

ID: 83821 Brandt 92000.00

ID: 98345 Kim 80000.00

Security issue – SQL Injection

(to be covered later)

- ❑ Basically, communicating with the database goes like this
 - Build a query string like
 - "select * from instructor where name = 'Srinivasan'"
 - from
 - The string "select * from instructor where name = '<input goes here>'"
 - and a value like Srinivasan received as input
 - Send this to the database and get the result in return
- ❑ One important security issue, to be covered later: **SQL Injection**
 - To interfere with the construction of the query string
 - example input to do injection:
 - X' or 'Y' = 'Y
 - a trick to get all data from all instructors

Functions and Procedures

Functions and Procedures

- ❑ Functions and procedures allow “**business logic**” to be **stored in the database** and executed from SQL statements as well as called from application programs.
- ❑ These can be defined either by the procedural component of SQL or by an external programming language such as Java, C, Python, R, and others.
- ❑ The syntax presented in the DB book is defined by the SQL standard.
 - Most databases implement nonstandard versions of this syntax.

Procedural Extensions and Stored Procedures

- ❑ SQL provides a **module language**
 - Permits definition of functions and procedures in SQL
- ❑ Functions
 - write your own functions and add them to the database
 - use them like any function predefined by the DBMS, that is, within SQL expressions
- ❑ Stored Procedures
 - store procedures in the database
 - execute them by "calling" them from applications or interfaces to the DBMS
 - this permits external applications to operate on the database without knowing about internal details
 - you can, for instance, develop a dedicated API that provides functionality but hides the database structure
- ❑ Triggers
 - you can add special procedures that are executed automatically by the system as a side effect of a modification to the database

Procedural Extensions and Stored Procedures

❑ PostgreSQL specialities

- PostgreSQL provides one of the most advanced frameworks and language extensions for adding functions and procedures to the DBMS

❑ PL/pgSQL

- PL/pgSQL is the PostgreSQL version of a **module language** (a procedural programming language)
- to be used for developing functions, procedures and triggers
- PL/pgSQL was originally inspired by Oracle's PL/SQL language.

Functions and Procedures

- ❑ Since SQL:1999 the standard supports functions and procedures
 - Functions/procedures can be written in SQL itself, or in an external programming language.
 - Some database systems (including PostgreSQL) support a particularly useful construct:
 - **table-valued function**, that returns a table/relation as a result
 - just as any SQL query does.
- ❑ SQL:1999 also supports a rich set of imperative constructs, including
 - Loops, if-then-else, assignment, and others
- ❑ Many database systems have proprietary procedural extensions to SQL that differ from the standard.

SQL Functions

❑ Define a function.

```
create function hello (s char(20))  
returns char(50)  
begin  
return concat('hello, ',s,'!');  
end;
```

```
create function hello (s char(20))  
returns char(50)  
language plpgsql as  
$$  
begin  
return concat('hello, ',s,'!');  
end;  
$$;
```

Why is \$\$ needed?

- “.” is the end-of-statement symbol
- not enclosing the function body in \$\$ would make “.” ambiguous

\$\$ is used to enclose the body as a literal

❑ Use the function.

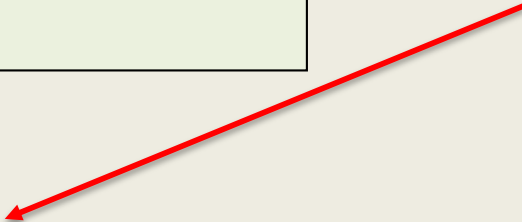
```
select hello('world') 'Message to all';
```

```
uni=# select hello('world') as "Message to all";  
Message to all  
-----  
hello, world!  
(1 row)
```

SQL Functions

```
create function hello (s char(20))
returns char(50)
language plpgsql as
$$
begin
return concat('hello, ',s,'!');
end;
$$;
```

Same function but
now used on a table



```
uni=# select hello(name) "Message to all" from instructor;
      Message to all
-----
hello, Srinivasan!
hello, Wu!
hello, Mozart!
hello, Einstein!
hello, El Said!
hello, Gold!
hello, Katz!
hello, Califieri!
hello, Singh!
```


SQL Functions

- ❑ Define a function that, given the name of a department, returns the **count of the number of instructors in that department**.

```
create function dept_count (dept_name varchar(20))  
returns integer  
begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
end
```

DSC Figure 5.6

- ❑ Find the department name and budget of all departments with more than 1 instructors.

```
select dept_name, budget  
from department  
where dept_count (dept_name) > 1
```

SQL Functions

- ❑ Same function, but now using the PL/pgSQL language
- ❑ Again **count of the number of instructors in that department.**

```
create function dept_count (d_name char(20))
returns integer
language plpgsql as $$
declare d_count integer;
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
    return d_count;
end;
$$;
```

- ❑ You can **call the dept_count()-function to get a value:**

```
uni=# select dept_count('Physics');
dept_count
-----
          2
(1 row)
```

- ❑ or use the dept_count()-function in a where-condition.

```
create function dept_count (d_name char(20))
returns integer
language plpgsql as $$
declare d_count integer;
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
    return d_count;
end;
$$;
```

```
uni=# select dept_name, budget from department
uni=# where dept_count(dept_name ) > 1;
 dept_name | budget
-----+-----
Comp. Sci. | 100000.00
Finance    | 120000.00
History    |  50000.00
Physics    |  70000.00
(4 rows)
```

- ❑ or use the dept_count()-function in the select clause

```
create function dept_count (d_name char(20))
returns integer
language plpgsql as $$
declare d_count integer;
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
    return d_count;
end;
$$;
```

```
uni=# select distinct dept_name, dept_count(dept_name)
uni-# from department;
```

dept_name	dept_count
Biology	1
Comp. Sci.	3
Elec. Eng.	1
Finance	2
History	2
Music	1

Standard (DSC book) vs PL/pgSQL

```
create function dept_count (dept_name varchar(20))  
returns integer  
begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
end
```

DSC Figure 5.6

Declarations outside
block

```
create function dept_count (d_name char(20))  
returns integer  
language plpgsql as $$  
declare d_count integer;  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = d_name;  
    return d_count;  
end;  
$$;
```

SQL Procedures

- ❑ The *dept_count* function could instead be written as a procedure:

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
```

```
    from instructor
```

```
    where instructor.dept_name = dept_name;
```

```
end
```

DSC page 200

- ❑ Procedures can be called from
 - other procedures or
 - SQL embedded in application programs or
 - command line, using the **call** statement.

SQL Procedures

- ❑ The *dept_count* function could instead be written as a procedure:

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
```

```
    from instructor
```

```
    where instructor.dept_name = dept_name;
```

```
end
```

DSC page 200

```
create procedure dept_count_proc (in d_name varchar(20),  
                                out d_count integer)  
  
language plpgsql  
as $$  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = d_name;  
end;  
$$;
```

SQL Procedures

```
create procedure dept_count_proc (in d_name varchar(20),  
                                out d_count integer)  
  
language plpgsql  
as $$  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = d_name;  
end;  
$$;
```

```
DO $$  
declare  
    n integer;  
begin  
    call dept_count_proc ('Comp. Sci.', n);  
    RAISE NOTICE 'Comp. Sci. has % instructors', n;  
end $$;
```

calling the procedure from an anonymous code block (DO-block)

```
NOTICE:  Comp. Sci. has 3 instructors  
DO  
uni=#
```

Output from RAISE NOTICE

- ❑ Nothing is gained from using the procedure here
A more straightforward coding would be a function as before:

```
create function dept_count_func (d_name varchar(20))  
returns integer  
language plpgsql as  
$$  
begin  
    return (select count(*)  
            from instructor  
            where instructor.dept_name = d_name);  
end;  
$$;
```

PL/pgSQL Function

As before, with
minor change

```
create procedure dept_count_proc (in d_name varchar(20),  
                                  out d_count integer)  
language plpgsql  
as $$  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = d_name;  
end;  
$$;
```

PL/pgSQL Procedure

SQL Functions

```
create function dept_count_func (d_name varchar(20))
returns integer
language plpgsql as
$$
begin
    return (select count(*)
            from instructor
            where instructor.dept_name = d_name);
end;
$$;
```

PL/pgSQL Function

❑ or an even simpler version as a plain SQL language function

```
create function dept_count_func (d_name varchar(20))
returns integer
language sql as
$$
    select count(*)
    from instructor
    where instructor.dept_name = d_name;
$$;
```

SQL Function

SQL Procedures

- ❑ So, why at all consider stored procedures?
- ❑ One important argument:
 - user-defined functions cannot execute transactions
 - stored procedures can do this
- ❑ Example: Transferring an amount from one account to another

```
drop table if exists accounts;
create table accounts (
    id int generated by default as identity,
    name varchar(100) not null,
    balance dec(15,2) not null,
    primary key(id)
);
insert into accounts(name,balance) values('Bob',10000);
insert into accounts(name,balance) values('Alice',10000);
```

```
uni=# select * from accounts;
 id | name  | balance
-----+-----+-----
  1 | Bob   | 10000.00
  2 | Alice | 10000.00
(2 rows)
```

SQL Procedures

- ❑ So, why at all consider stored procedures?
- ❑ One important argument:
 - user-defined functions cannot execute transactions
 - stored procedures can do this
- ❑ Example: Transferring amount from one account to another

```
drop table if exists accounts;
```

```
create procedure transfer(sender int, receiver int, amount dec)
language plpgsql
as $$
begin
    -- subtracting the amount from the sender's account
    update accounts
    set balance = balance - amount
    where id = sender;
    -- adding the amount to the receiver's account
    update accounts
    set balance = balance + amount
    where id = receiver;
    commit;
end; $$
```

SQL Procedures

- ❑ So, why at all consider stored procedures?
- ❑ One important argument:
 - user-defined functions cannot execute transactions
 - stored procedures can do this
- ❑ Example: Transferring amount from one account to another

```
drop table if exists accounts;
```

```
create procedure transfer(sender int, receiver int, amount dec)
language plpgsql
as $$
begin
    -- subtracting the amount from the sender's account
    update accounts
    set balance = balance - amount
    where id = sender;
    -- adding the amount to the receiver's account
    update accounts
    set balance = balance + amount
    where id = receiver;
    commit;
end;$$
```

```
uni=# call transfer(1,2,500);
```

```
uni=# select * from accounts;
 id | name  | balance
----+-----+-----
  1 | Bob   |  9500.00
  2 | Alice | 10500.00
(2 rows)
```

Table Functions

- ❑ The SQL standard supports functions that can return tables as results; such functions are called **table functions**

- ❑ Example: Return all instructors in a given department

create function *instructor_of* (*dept_name* **char**(20))

returns table (

ID **varchar**(5),
name **varchar**(20),
dept_name **varchar**(20),
salary **numeric**(8,2))

return table

(**select** *ID, name, dept_name, salary*
from *instructor*
where *instructor.dept_name = instructor_of.dept_name*);

DSC Figure 5.7

- ❑ Usage

select *
from table (*instructor_of* ('Music'))

Table Functions

- ❑ Same function, here using plain SQL (in the body)

```
create or replace function instructors_of (dept_name char(20))
    returns table (ID varchar(5),
                  name varchar(20),
                  dept_name varchar(20),
                  salary numeric(8,2))
language sql as
$$
    select ID, name, dept_name, salary
    from instructor
    where instructor.dept_name = instructors_of.dept_name;
$$;
```

- ❑ Here used to retrieve instructors in the Physics department.

```
uni=# select * from instructors_of ('Physics');
 id   | name   | dept_name | salary
-----+-----+-----+-----
 22222 | Einstein | Physics   | 95000.00
 33456 | Gold    | Physics   | 87000.00
(2 rows)
```

Language Constructs for Procedures & Functions

- ❑ SQL supports constructs that gives it almost all the power of a general-purpose programming language.
- ❑ Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statements
- ❑ Loops using (among other) while and repeat statements:
 - **while** boolean expression **do**
 sequence of statements ;
end while
 - **repeat**
 sequence of statements ;
 until boolean expression
end repeat

Language Constructs – if-then-else

- ❑ Conditional statements (**if-then-else**)
 - if** *boolean expression*
 - then** *statement or compound statement*
 - elseif** *boolean expression*
 - then** *statement or compound statement*
 - else** *statement or compound statement*
 - end if**

SQL Procedure, example with IF and WHILE loop

```
drop table if exists foo;  
create table foo  
(  
    id serial primary key,  
    val integer  
);
```

a table, **foo**, for testing

defining the **load_foo()** procedure

calling and showing the effect

```
create or replace procedure load_foo()  
language plpgsql as  
$$  
declare  
    i_max integer := 4;  
    i integer := 0;  
    n integer;  
begin  
    while i < i_max loop  
        n:=(random() * 10000);  
        if n>5000 then n=0;  
        end if;  
        insert into foo (val) values (n);  
        i:=i+1;  
    end loop;  
end  
$$;
```

```
uni=# call load_foo();  
CALL  
  
uni=# select * from foo;  
   id | val  
-----+-----  
    9 |    0  
   10 |    0  
   11 | 4978  
   12 |    0  
(4 rows)
```

SQL Procedure, example with IF and WHILE loop

❑ Notice SQL-details

- drop ... if exists ... (very useful in a script you want to run repeatedly)
 - `drop table if exists foo;`
- auto incrementing primary key
 - `id serial primary key,`
- declaration and initialization of variable
 - `i_max integer := 4;`
- while loop to do several DML-statements
 - `while i < i_max loop`
- random() between 0 and 1 used to generate number between 0 and 9999
 - `n:=(random() * 10000);`
- if statement to replace numbers > 5000 with 0
 - `if n>5000 then n=0;`

Calling a procedure from another

```
drop table if exists foo;  
create table foo  
(  
    id serial primary key,  
    val integer  
);
```

a table, **foo**, for testing

calling and showing the effect

```
uni=# call test();  
CALL
```

```
uni=# select * from foo;  
 id | val  
----+-----  
 13 |    0  
 14 | 2068  
 15 | 4190  
 16 | 2770  
 17 |    0  
 18 |   211  
 19 | 2982  
 20 |    0  
(8 rows)
```

defining the procedure that calls the procedure

```
create or replace procedure test()  
language plpgsql as  
$$  
begin  
    call load_foo();  
    call load_foo();  
end  
$$;
```

SQL Procedure, same functionality, but simplified

```
create or replace procedure load_foo()  
language plpgsql as  
$$  
declare  
    i_max integer := 4;  
    i integer := 0;  
    n integer;  
begin  
    while i < i_max loop  
        n:=(random() * 10000);  
        if n>5000 then n=0;  
        end if;  
        insert into foo (val) values (n);  
        i:=i+1;  
    end loop;  
end  
$$;
```

```
create or replace procedure load_foo()  
language plpgsql as  
$$  
begin  
    for i in 1..4 loop  
        insert into foo (val) values (random() * 10000);  
    end loop;  
end  
$$;
```

Cursor

- declared by a query
- supports row-by-row traversal of the result of the query

❑ **declare**

- Before a cursor can be used it must be declared (defined).
- declare `curl cursor for select name,salary from instructor;`

❑ **open** – perform the query

- The cursor must be opened for use. This process actually retrieves the data using the previously defined SELECT statement.
- `open curl;`

❑ **fetch** – get the next row from the table

- Individual rows can be fetched (retrieved) as needed.
- `fetch curl into ...;`

❑ **close** – close the cursor (clean up)

- When done, the cursor must be closed.
- `close curl;`

SQL Procedure using cursor, example

a table, **vip**, for testing

```
drop table if exists vip;
create table vip as
  select name, salary
  from instructor;
truncate vip;
```

```
create or replace procedure curdemo()
language plpgsql as
$$
declare
  rec record;
  curl cursor for select name,salary from instructor;
begin
  open curl;
  loop
    fetch curl into rec;
    exit when not found;
    if rec.salary > 81000 then
      insert into vip
        values (rec.name,rec.salary);
    end if;
  end loop;
  close curl;
end;
$$;
```

calling and showing the effect

```
uni=# call curdemo();
CALL

uni=# select * from vip;
   name  | salary
-----+-----
 Wu      | 90000.00
 Einstein | 95000.00
 Gold    | 87000.00
 Brandt  | 92000.00
(4 rows)
```

SQL Procedure using cursor, example(cont.)

❑ Notice SQL-details

- the four “using cursors”-issues to remember:
 - **declare, open, fetch, close**
- conditional statement (fairly standard)
 - `if ... then ... end if;`
- a very useful data type **record**:
 - `rec record;`
 - `...`
 - `fetch cur1 into rec;`
- another loop construction
 - `loop ... exit when ... end loop;`

Yet another loop ... to replace a normal cursor

- ❑ A very useful loop in PostgreSQL is the following

```
[ <<label>> ]  
FOR target IN query LOOP  
    statements  
END LOOP [ label ];
```

- ❑ Testing here with a **DO-block**
 - anonymous function
 - can be used for adhoc tasks and for testing expressions
- ❑ raise notice ...
 - a kind of print() statement
 - can be used while testing

the FOR ... IN loop provides a kind of "implicit" cursor

testing with a DO block

```
uni=#  
do $$  
    declare  
        rec record;  
    begin  
        for rec in select name  
            from instructor  
        loop  
            raise notice '%', rec.name;  
        end loop;  
    end;  
$$;  
NOTICE:  Srinivasan  
NOTICE:  Wu  
NOTICE:  Mozart  
NOTICE:  Einstein  
NOTICE:  El Said  
NOTICE:  Gold  
NOTICE:  Katz  
NOTICE:  Califieri
```

SQL using cursor (now implicit), example

a table, **vip**, for testing

-- vip(name, salary)
-- see slide 43

```
create or replace procedure curdemo2() as
$$
declare
    rec record;
begin
    for rec in select name,salary from instructor
    loop
        if rec.salary > 81000 then
            insert into vip
                values (rec.name,rec.salary);
        end if;
    end loop;
end;
$$
language plpgsql;
```

calling and showing the effect

```
uni=# truncate vip;
TRUNCATE TABLE
uni=# call curdemo2();
CALL
uni=# select * from vip;
      name      | salary
-----+-----
Wu              | 90000.00
Einstein        | 95000.00
Gold            | 87000.00
Brandt          | 92000.00
(4 rows)
```

SQL using cursor (now implicit), ex. (cont.)

- ❑ Compare to the cursor examples above (slide 43 and slide 46)
 - the loop is changed to
 - **for rec in** `select name,salary from instructor`
 - **loop**
 - ...
 - **end loop**
 - the cursor is replaced by the expression given as argument in the for loop
 - `for rec in` **select** `name,salary` **from** `instructor`
 - conceptually this is still a cursor

External Language Routines

- ❑ SQL allows us to define functions in a programming language such as Java, C, Python, R, and others.
 - Can be more efficient than functions defined in SQL, and computations that cannot be carried out in SQL can be executed by these functions.
- ❑ Declaring external language procedures and functions
Examples:

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out count integer)
```

```
language C
```

```
external name '/usr/avi/bin/dept_count_proc'
```

```
create function dept_count(dept_name varchar(20))
```

```
returns integer
```

```
language C
```

```
external name '/usr/avi/bin/dept_count'
```

External Language Functions/Procedures

- ❑ Notice the PostgreSQL **CREATE FUNCTION** statement:

```
CREATE FUNCTION function_name(...)
RETURNS type AS
$$
BEGIN
-- logic
END;
$$
LANGUAGE language_name;
```

- ❑ By default, PostgreSQL supports **5 languages**:
 - SQL, PL/pgSQL, PL/Tcl, PL/Perl and PL/Python
- ❑ You can install other procedural languages
 - e.g., Java, Lua, R, Unix Shell, JavaScript, C++, ...

External Language Routines (Cont.)

- ❑ Benefits of external language functions/procedures:
 - more efficient for many operations, and more expressive power.
- ❑ Drawbacks
 - Code to implement function may need to be loaded into database system and executed in the database system's address space.
 - risk of accidental corruption of database structures
 - security risk, allowing users access to unauthorized data
 - There are alternatives, which give good security at the cost of potentially worse performance.

Why use Stored functions and procedures?

- ❑ Stored functions and procedures (routines) can be particularly useful
 - When **multiple client applications** are written in different languages or work **on different platforms**, but **need to perform the same database operations**.
 - When security is paramount. **Banks**, for example, **use stored procedures** and functions for all common operations
 - In addition, you can store **libraries of functions and procedures** in the database server
 - Provide **improved performance**. Less information needs to be sent between the server and the client.
 - Tradeoff: **increase the load on the database server**

Triggers

Triggers

- ❑ A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- ❑ To design a trigger, we must:
 - Specify the **conditions** under which the trigger is to be executed.
 - Specify the **actions** to be taken when the trigger executes.
- ❑ Triggers were introduced to the SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.

Triggering Events and Actions in SQL

- ❑ Triggering event can be **insert**, **delete** or **update**
- ❑ Triggers on update can be restricted to specific attributes
 - For example, **after update of** *takes on grade*
- ❑ Values of attributes before and after an update can be referenced
 - **referencing old row as** : for deletes and updates
 - **referencing new row as** : for inserts and updates
- ❑ Triggers can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.

```
create trigger setnull_trigger before update of takes  
referencing new row as nrow  
for each row  
  when (nrow.grade = ' ')  
  begin atomic  
    set nrow.grade = null;  
  end;
```

PostgreSQL Triggers

- ❑ To create a new trigger in PostgreSQL:
 - Create a trigger function using CREATE FUNCTION statement.
 - Bind this trigger function to a table using CREATE TRIGGER statement.
- ❑ Create the **trigger function**

```
CREATE FUNCTION trigger_function() RETURN trigger AS
```

- a function similar to an ordinary function,
- does not take any arguments
- has return **return type trigger**
- important variables: **OLD** and **NEW** represent the states of row in the table before or after the triggering event.

PostgreSQL Triggers

- ❑ Create the trigger
 - use the **CREATE TRIGGER** statement:

```
CREATE TRIGGER trigger_name {BEFORE | AFTER | INSTEAD OF} {event [OR ...]}  
ON table_name  
[FOR [EACH] {ROW | STATEMENT}]  
EXECUTE PROCEDURE trigger_function
```

- The event could be INSERT, UPDATE, DELETE or TRUNCATE.
- BEFORE or AFTER event specifies the order of the trigger and the update
- INSTEAD OF is used only for views
- two kinds of triggers: row level trigger and statement level trigger,

Trigger Example – Referential constraint

- ❑ E.g. *time_slot_id* is not a primary key of *timeslot*, so we cannot create a foreign key constraint from *section* to *timeslot*.
- ❑ Alternative: use triggers on *section* and *timeslot* to enforce integrity constraints
- ❑ Figure 5.9 in DSC book:

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from timeslot)) /* time_slot_id not present in timeslot */
begin
    rollback
end;
```

Will not work in PostgreSQL

Trigger Example – Referential constraint

❑ Figure 5.9 in DSC book does NOT work

- Rollback is not allowed in a trigger in PostgreSQL

❑ The following is an alternative

- The result is the same: an update with a time_slot_id not present in the time_slot table will not be allowed (and will thus be ignored)

❑ Figure 5.9 in DSC book:

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from time_slot)) /* time_slot_id not present in time_slot */
begin
    rollback
end;
```

Will not work in PostgreSQL

```
create function timecheck() -- the trigger function
returns trigger as $$
begin
    if (new.time_slot_id not in (select time_slot_id from time_slot)) then
        raise exception 'time_slot_id is unknown';
    end if;
    return new;
end; $$
language plpgsql;

create trigger timecheck_trig -- the trigger (calling the trigger function)
before insert on section
for each row execute procedure timecheck();
```

```
uni=# insert into section values
uni-# ('BIO-301', '2', 'Winter', '2009', 'Painter', '514', 'I');
ERROR:  time_slot_id is unknown
CONTEXT:  PL/pgSQL function timecheck() line 4 at RAISE
```

❑ Notice SQL and PostgreSQL details

- The example is a **before** rather than an **after** trigger
- **if** (inside the block) replaces **when** (outside)
- “**referencing new row as *nrow***” won’t work, but you can reference the new value simply with **new**
 - **new** can be used in **insert** and **update**-triggers
 - **old** can be used similarly in **delete** and **update**-triggers
- **raise exception** will prevent the insert and return an error message

```
create function timecheck() -- the trigger function
returns trigger as $$
begin
    if (new.time_slot_id not in (select time_slot_id from time_slot)) then
        raise exception 'time_slot_id is unknown';
    end if;
    return new;
end; $$
language plpgsql;
```

```
create trigger timecheck_trig -- the trigger (calling the trigger function)
before insert on section
for each row execute procedure timecheck();
```

```
uni=# insert into section values
uni-# ('BIO-301', '2', 'Winter', '2009', 'Painter', '514', 'I');
ERROR:  time_slot_id is unknown
CONTEXT:  PL/pgSQL function timecheck() line 4 at RAISE
```

Trigger Example – Ad hoc constraint

- ❑ Company policy (insert on instructor trigger)
 - No new employments in high budget departments (≥ 90000)
 - New employees (instructors) must never have a salary greater than everybody else

```
drop trigger if exists instructorcheck_trig on section;
drop function if exists instructorcheck;
create function instructorcheck() -- the trigger function
returns trigger as $$
begin
    if (new.dept_name not in (select dept_name from department where budget < 90000)) then
        raise exception 'No no no, no new employees in the % department', new.dept_name;
    end if;
    if (new.salary > (select max(salary) from instructor)) then
        raise exception 'No no no, salary too high';
    end if;
    return new;
end; $$
language plpgsql;
```

% is a placeholder to be replaced by the value of **new.dept_name**

```
create trigger instructorcheck_trig -- the trigger
before insert on instructor for each row execute procedure instructorcheck();
```

```
uni=# insert into instructor values (12345, 'Wong', 'Finance', 80000);
ERROR:  No no no, no new employees in the Finance department
```

```
uni=# insert into instructor values (23456, 'Wang', 'History', 100000);
ERROR:  No no no, salary too high
```


Triggering Events and Actions in SQL

- ❑ Triggering event can be **insert**, **delete** or **update**
- ❑ Triggers can be activated before an event, which can serve as extra constraints. E.g. convert blank grades to null.
- ❑ Triggers can also be activated after an event

```
create function setnull() -- the trigger function
returns trigger as $$
begin
    if (new.grade <= ' ') then
        new.grade := null;
    end if;
    RETURN NEW;
end;$$
language plpgsql;

create trigger setnull_trig -- the trigger
before update on takes for each row execute procedure
setnull();
```

T

```
create function setnull() -- the trigger function
returns trigger as $$
begin
    if (new.grade = ' ') then
        new.grade := null;
    end if;
    RETURN NEW;
end;$$
language plpgsql;

create trigger setnull_trig -- the trigger
before update on takes for each row execute procedure
setnull();
```

testing ...

```
uni=# select * from takes where grade is null;
```

id	course_id	sec_id	semester	year	grade
98988	BIO-301	1	Summer	2018	

(1 row)

```
uni=# update takes
```

```
uni=# set grade=' ' where id ='98765' and course_id='CS-101';
UPDATE 1
```

```
uni=# select * from takes where grade is null;
```

id	course_id	sec_id	semester	year	grade
98988	BIO-301	1	Summer	2018	
98765	CS-101	1	Fall	2017	

(2 rows)

Trigger example SKAL RETTES

- ❑ **An update** trigger ensuring that amount on account always satisfies $0 \leq \text{amount} \leq 100$

```
create function upd_check() before update on account
returns trigger as $$
begin
    if new.amount < 0 then
        set new.amount = 0;
    elseif new.amount > 100 then
        set new.amount = 100;
    end if;
    return new;
end;

create trigger upd_check_trig -- the trigger
after update on takes for each row execute procedure
upd_check();
```

Trigger to update tot_cred on student

❑ Keeping the value of tot_cred up-to-date

```
create trigger credits_earned after update of takes on grade
referencing new row as nrow
referencing old row as orow
for each row
when nrow.grade <> 'F' and nrow.grade is not null
    and (orow.grade = 'F' or orow.grade is null)
begin atomic
    update student
    set tot_cred = tot_cred +
        (select credits
         from course
         where course.course_id = nrow.course_id)
    where student.id = nrow.id;
end;
```

DSC Figure 5.10

Trigger to update tot_cred on student

- ❑ Keeping the value of tot_cred up-to-date
- ❑ In PostgreSQL with
 - a trigger function and
 - a trigger

```
create function update_tot_cred() -- the trigger function
returns trigger as $$
begin
    update student
    set tot_cred=(select sum(credits)
                  from takes join course using (course_id)
                  where takes.id=student.id)
    where student.id=new.id;
    return NEW;
end; $$
language plpgsql;

create trigger update_tot_cred_trig -- the trigger
after insert on takes
for each row execute procedure update_tot_cred();
```

Trigger to update tot_cred on student

- ❑ Keeping the value of tot_cred up-to-date

- ❑ In PostgreSQL with

- a trigger function and
- a trigger

- ❑ Testing

```
uni=# select * from student where id='19991';
```

id	name	dept_name	tot_cred
19991	Brandt	History	80

(1 row)

```
uni=# insert into takes values ('19991', 'CS-190', '2', 'Spring', '2017', 'A');  
INSERT 0 1
```

```
uni=# select * from student where id='19991';
```

id	name	dept_name	tot_cred
19991	Brandt	History	7

(1 row)

Statement Level Triggers

- ❑ Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction.
 - Can be more efficient when dealing with SQL statements that update a large number of rows
- ❑ **Supported by** some DBMS' (including **Postgres**) using
 - Use **for each statement** instead of **for each row**
 - Use **referencing old table** or **referencing new table** to refer to temporary tables (called *transition tables*) containing the affected rows
- ❑ Insertion of 887000 rows:
insert into movie.movie
select id, title, production_year
from imdb_movie.movie where kind_id=1;
- ❑ with row-level: 887000 actions, with statement level: 1 action

When Not To Use Triggers

- ❑ Triggers were used earlier for tasks such as
 - Maintaining **summary data** (e.g., total salary of each department)
 - **Replicating databases** by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- ❑ There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication
- ❑ Encapsulation can be used instead of triggers in many cases
 - Define methods og SQL procedures to update fields
 - Carry out actions as part of these