

The main purpose of this assignment is to practice querying using SQL. The target is a larger university database with 2000 students and around 28.500 course enrollments (the takes relation). You can build this database on your local PostgreSQL server by running a script **university\_large.sql**. Check out the instructions by the end of this text on: building the database, producing an output file and handing in.

[illegible]

Notice that the primary keys are shown (underlined) and the foreign keys are indicated by arrows. Some of the primary keys (and thereby also some of the foreign keys) are composite.

Try to study the content of the database before you begin. Run some example queries and check the type of values used for e.g. keys.

Write SQL queries to answer the questions below. Notice that query-answers are listed, so for each you will “simply” have to find an SQL-expression, that produces the given answer.  
See instructions, at the end of this note.

You are supposed to submit one hand in per group and we strongly recommend you to work with the questions individually as well as in groups. Discuss and decide on what you consider to be the best solutions and hand these in.

advisor	2000
classroom	30
course	200
department	20
instructor	50
prereq	100
section	103
student	2000
takes	28544
teaches	103
time_slot	20

## Questions

**Observe:** Most of the queries to answer the questions below should return almost immediately. Stop your query evaluation if you experience long waiting and reconsider your query expression. You can skip those questions that you cannot find an answer to.

-- 1. Find the names of all the instructors from Biology department

```
name
-----
Queiroz
Valtchev
(2 rows)
Time: 2.481 ms
```

-- 2. Find the names of courses in Computer science department which have 3 credits

```
title
-----
International Finance
Japanese
Computability Theory
(3 rows)
Time: 1.146 ms
```

-- 3. For the student with ID 30397, show all course\_id and title of all courses registered for by the student.

```
course_id | title
-----+-----
105       | Image Processing
158       | Elastic Structures
200       | The Music of the Ramones
319       | World History
349       | Networking
461       | Physical Chemistry
468       | Fractal Geometry
496       | Aquatic Chemistry
626       | Multimedia Design
631       | Plasma Physics
642       | Video Gaming
702       | Arabic
760       | How to Groom your Cat
795       | Death and Taxes
959       | Bacteriology
960       | Tort Law
(16 rows)
Time: 1.352 ms
```

-- 4. As above, but show the total number of credits for such courses (taken by that student).

-- Don't display the tot\_creds value from the student table, you should use SQL aggregation on courses taken by the student.

```
course_id | title | sum
-----+-----+-----
105       | Image Processing | 6
158       | Elastic Structures | 3
200       | The Music of the Ramones | 4
319       | World History | 4
349       | Networking | 4
461       | Physical Chemistry | 3
468       | Fractal Geometry | 8
496       | Aquatic Chemistry | 3
626       | Multimedia Design | 4
631       | Plasma Physics | 4
642       | Video Gaming | 3
```

## Complex IT Systems

702	Arabic	3
760	How to Groom your Cat	3
795	Death and Taxes	3
959	Bacteriology	4
960	Tort Law	3

(16 rows)

Time: 0.468 ms

-- 5. Now display the total credits (over all courses) for each of the students having more than 85 in total credits, along with the ID of the student;

-- don't bother about the name of the student.

-- don't bother about students who have not registered

id	sum
72669	87
50039	86
47379	86
107	86
44551	87
12078	93

(6 rows)

Time: 36.310 ms

-- 6. Find the names of all students who have taken any course at the Languages department with the grade 'A+' (there should be no duplicate names)

name
Abraham
Bhattacharya
Boldin
Carey
Cheah
Ching
Cochran
Damas
Denso
Ebou
Frangeu
Geißl
Hughes
Januszcwski
Kirtane
Komori
Kwan
Macias
Masri
Nakajima
Oki
Oller
Palomo
Paniez
Patne
Pavlovico
Planti
Roses
Saill
Savolainen
Thimm
Vries
Wingb
Wood
Zafar

(35 rows)

Time: 7.196 ms

-- 7. Display the IDs of all instructors from the Marketing department who have never taught a course (interpret "taught" as "taught or is scheduled to teach")

```
id
-----
58558
96895
74426
(3 rows)
Time: 1.017 ms
```

-- 8. As above, but display the names of the instructors also, not just the IDs.

```
id | name
-----+-----
96895 | Mird
74426 | Kenje
58558 | Dusserre
(3 rows)
Time: 0.474 ms
```

-- 9. Using the university schema, write an SQL query to find the number of students in each section in year 2009.

-- The result columns should be "course\_id, sec\_id, year, semester, num", where the latter is the number.

-- You do not need to output sections with 0 students.

```
course_id | sec_id | semester | year | num
-----+-----+-----+-----+-----
972       | 1      | Spring   | 2009 | 280
960       | 1      | Fall     | 2009 | 307
304       | 1      | Fall     | 2009 | 307
604       | 1      | Spring   | 2009 | 300
105       | 1      | Fall     | 2009 | 327
334       | 1      | Fall     | 2009 | 268
237       | 2      | Fall     | 2009 | 311
486       | 1      | Fall     | 2009 | 304
(8 rows)
Time: 7.475 ms
```

-- 10. Find the maximum and minimum enrollment across all sections, considering only sections that had some enrollment, don't worry about those that had no students taking that section.

-- 10a. using a subquery in from

```
max | min
-----+-----
338 | 264
(1 row)
Time: 36.239 ms
```

-- 10b. using a with statement

```
max | min
-----+-----
338 | 264
(1 row)
Time: 19.227 ms
```

-- 11. Find all sections that had the maximum enrollment (along with the enrollment),

-- 11a. using a subquery.

```
course_id | sec_id | semester | year | num
-----+-----+-----+-----+-----
362       | 1      | Fall     | 2005 | 338
192       | 1      | Fall     | 2002 | 338
(2 rows)
Time: 74.180 ms
```

-- 11b. using a with clause (e.g. defining temporary enrollment and maxnumber tables)

```
course_id | sec_id | year | semester | num
-----+-----+-----+-----+-----
362       | 1      | 2005 | Fall     | 338
192       | 1      | 2002 | Fall     | 338
(2 rows)
Time: 19.809 ms
```

-- 12. As in in Q10, but now also include sections with no students taking them; the enrollment for such sections should be treated as 0.

-- Use aggregation and outer join -- use aggregation on a left outer join

```
max | min
-----+-----
338 | 0
(1 row)
Time: 56.797 ms
```

-- 13. Find all courses that the instructor with id '19368' have taught

```
id | course_id | sec_id | semester | year
-----+-----+-----+-----+-----
19368 | 581       | 1      | Spring   | 2005
19368 | 545       | 1      | Fall     | 2001
19368 | 591       | 1      | Spring   | 2005
(3 rows)
Time: 0.469 ms
```

-- 14. Find instructors who have taught all the above courses  
-- (use "... not exists ... except ...")

```
id
-----
19368
41930
(2 rows)
Time: 20.633 ms
```

-- 15. Insert each instructor as a student, with tot\_creds = 0, in the same department

```
INSERT 0 47
Time: 7.700 ms
```

-- 16. Now delete all the newly added "students" above (note: already existing students who happened to have tot\_creds = 0 should not get deleted)

```
DELETE 47
Time: 5.191 ms
```

-- 17. Some of you may have noticed that the tot\_cred value for students do not always match the credits from courses they have taken.

```
-- Write a query to compare these that show a students id and
tot_cred together with the correct calculated total credits.
-- Show only the cases where the two values match
```

```
id | tot_cred | sum
-----+-----+-----
30177 | 41 | 41
38336 | 39 | 39
39520 | 43 | 43
14094 | 45 | 45
8378 | 47 | 47
48901 | 57 | 57
54508 | 41 | 41
41596 | 51 | 51
44584 | 58 | 58
16907 | 59 | 59
29920 | 62 | 62
61232 | 48 | 48
57787 | 60 | 60
67340 | 38 | 38
63390 | 41 | 41
(15 rows)
Time: 44.004 ms
```

```
-- 18. Write and execute a query to update tot_cred for all
students to the correct calculated value based on the credits passed
- thereby bringing the database back to consistency.
```

```
UPDATE 2000
Time: 266.306 ms
```

```
-- 19. Run Q17 again, but now only show when the two values differ
```

```
id | tot_cred | sum
----+-----+-----
(0 rows)
Time: 44.599 ms
```

```
-- 20. Update the salary of each instructor to 29001 + 10000
times the number of course sections they have taught.
```

```
UPDATE 50
Time: 1.791 ms
```

```
-- 21. List name and salary for the 10 first instructors (alphabetic
order)
```

```
id | name | salary
-----+-----+-----
37687 | Arias | 29001.00
95030 | Arinb | 29001.00
28400 | Atanassov | 49001.00
52647 | Bancelhon | 29001.00
15347 | Bawa | 39001.00
97302 | Bertolino | 29001.00
90376 | Bietzk | 39001.00
34175 | Bondi | 59001.00
3335 | Bourrier | 49001.00
90643 | Choll | 39001.00
(10 rows)
Time: 0.233 ms
```

## How to (re)build the larger university database

You can download and install the database on your own computer. Do the following:

- Download the SQL script file **university\_large.sql** from the CIT Moodle site.
- Open your command line interface and change to the directory where you placed the file **university\_large.sql**.

- Run the two commands  
**psql -U postgres -c "create database university\_large"**  
**psql -U postgres -d university\_large -f university\_large.sql**
- You may consider to drop the database if you need to start with a fresh database content while working or when you are finished using the database for Assignment 1. One way to drop the database is by using this command:  
**psql -U postgres -c "drop database university\_large"**

## How to produce an output file and hand in your solution

When you have tested all your solutions to the questions one by one, include all these in a single SQL script file **citXX-assignment1.sql** (where x is your group number) with content like:

```
-- GROUP: citXX, MEMBERS: <name1>, <name2>, ...
-- 1.
SELECT name from instructor where dept_name='Biology';

-- 2.
SELECT ...

...

-- 3.
SELECT ...

...

...
```

To indicate “no solution” for a question, just write “-- no solution” in place of the SELECT expression in the script. To include timings in the output (not required, not shown above) you can add a first line:

**\timing.** To hand in your solution do the following.

- Generate an output file **citXX-assignment1.txt** (where XX is your group number) by running the command:  
**psql -U postgres -a -d university\_large -f citXX-assignment1.sql > citXX-assignment1.txt**  
What’s going on here is that all the SQL-code in your input file **citXX-assignment1.sql** will be processed and the output will be written to the file **citXX-assignment1.txt** including echoes of the SQL expressions and the -- comments (the -a option takes care of the echoing). See all the available options for the psql command is [PGMAN] [psql command](#).
- Check your output file **citXX-assignment1.txt** and hand it in together with your script file **citXX-assignment1.sql** on the Moodle page.  
OBS: Just one hand-in per group, but remember to list all group members names in the beginning of the script file

Main reference for this assignment is:

- [DSC] Database System Concepts, Abraham Silberschatz, Henry Korth, S. Sudarshan, 7th Edition, 2019, chapter 3 and 4

To read more about how to use the command line tool **psql** consider:

- PostgreSQL Tutorial, 17 Practical psql Commands That You Don’t Want To Miss  
<http://www.postgresqltutorial.com/psql-commands/>
- [PGMAN] psql command  
<https://www.postgresql.org/docs/current/app-psql.html>